

ASSIGNMENT 6

Om Gupta - 22110174

Pranav Somase - 22110200

GITHUB - https://github.com/hiomgupta/STTAi_L6

Data Preprocessing

```
[66] # Load the dataset
iris = load_iris()
X, y = iris.data, iris.target.reshape(-1, 1)

# One-hot encoding of labels
encoder = OneHotEncoder(sparse_output=False)
y_onehot = encoder.fit_transform(y)

# Train-test split
X_train, X_temp, y_train, y_temp = train_test_split(X, y_onehot, test_size=0.3, random_state=42)
X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp, test_size=0.66, random_state=42)

# Normalize feature values
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_val = scaler.transform(X_val)
X_test = scaler.transform(X_test)

# Print sample training data
print("Sample Training Data:")
print(pd.DataFrame(X_train[:5], columns=['Feature1', 'Feature2', 'Feature3', 'Feature4']))
print("\nCorresponding Labels:")
print(np.argmax(y_train[:5], axis=1))

Sample Training Data:
   Feature1  Feature2  Feature3  Feature4
0 -0.413416 -1.462003 -0.099511 -0.323398
1  0.551222 -0.502563  0.717703  0.353032
2  0.671802  0.217016  0.051192  0.758090
3  0.912901 -0.822844  0.309096  0.217740
4  1.636440  1.416315  1.301427  1.705891

Corresponding Labels:
[1 2 2 1 2]
```

Model 1: Multi-Layer Perceptron (MLP)

```
# Initialize Weights & Biases
wandb.init(project="mlp-iris-experiment", config={"learning_rate": 0.001, "batch_size": 32, "epochs": 50})

# Define MLP model using Keras
model = tf.keras.Sequential([
    tf.keras.layers.Dense(16, activation="relu", input_shape=(4,)),
    tf.keras.layers.Dense(3, activation="softmax")
])

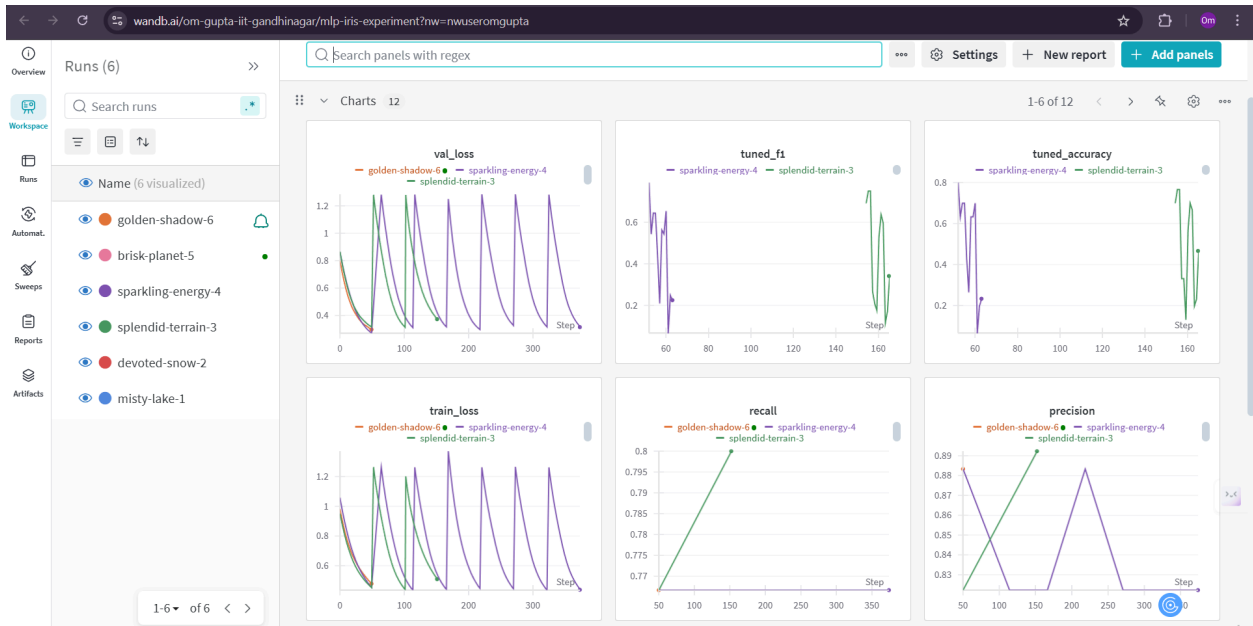
# Compile the model
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Train the model and log metrics
history = model.fit(X_train, y_train, validation_data=(X_val, y_val), epochs=50, batch_size=32)

# Log training loss & accuracy to W&B
for epoch in range(len(history.history['loss'])):
    wandb.log({
        "epoch": epoch,
        "train_loss": history.history['loss'][epoch],
        "val_loss": history.history['val_loss'][epoch]
    })

/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:47: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using an 'Input(shape)' object as the first layer in t
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/50
4/4 — 5s 469ms/step - accuracy: 0.5130 - loss: 1.2161 - val_accuracy: 0.4667 - val_loss: 1.2844
Epoch 2/50
4/4 — 0s 9ms/step - accuracy: 0.5651 - loss: 1.2165 - val_accuracy: 0.5333 - val_loss: 1.2463
Epoch 3/50
4/4 — 0s 9ms/step - accuracy: 0.6293 - loss: 1.1403 - val_accuracy: 0.5333 - val_loss: 1.2807
Epoch 4/50
4/4 — 0s 9ms/step - accuracy: 0.5605 - loss: 1.1689 - val_accuracy: 0.5333 - val_loss: 1.1732
```

Epoch 24/50	0s 9ms/step	- accuracy: 0.7756	- loss: 0.7006	- val_accuracy: 0.8000	- val_loss: 0.6151
4/4					
Epoch 25/50	0s 10ms/step	- accuracy: 0.8051	- loss: 0.6773	- val_accuracy: 0.8000	- val_loss: 0.5943
4/4					
Epoch 26/50	0s 9ms/step	- accuracy: 0.8138	- loss: 0.6694	- val_accuracy: 0.8000	- val_loss: 0.5741
4/4					
Epoch 27/50	0s 17ms/step	- accuracy: 0.8179	- loss: 0.6334	- val_accuracy: 0.8667	- val_loss: 0.5551
4/4					
Epoch 28/50	0s 10ms/step	- accuracy: 0.8144	- loss: 0.6551	- val_accuracy: 0.8667	- val_loss: 0.5374
4/4					
Epoch 29/50	0s 10ms/step	- accuracy: 0.8089	- loss: 0.6268	- val_accuracy: 0.8667	- val_loss: 0.5212
4/4					
Epoch 30/50	0s 9ms/step	- accuracy: 0.8457	- loss: 0.6243	- val_accuracy: 0.8667	- val_loss: 0.5048
4/4					
Epoch 31/50	0s 12ms/step	- accuracy: 0.8255	- loss: 0.6138	- val_accuracy: 0.8667	- val_loss: 0.4897
4/4					
Epoch 32/50	0s 12ms/step	- accuracy: 0.8238	- loss: 0.5811	- val_accuracy: 0.8667	- val_loss: 0.4749
4/4					
Epoch 33/50	0s 10ms/step	- accuracy: 0.8380	- loss: 0.5757	- val_accuracy: 0.9333	- val_loss: 0.4611
4/4					
Epoch 34/50	0s 9ms/step	- accuracy: 0.8418	- loss: 0.5766	- val_accuracy: 0.9333	- val_loss: 0.4492
4/4					
Epoch 35/50	0s 9ms/step	- accuracy: 0.8481	- loss: 0.5632	- val_accuracy: 0.9333	- val_loss: 0.4383
4/4					
Epoch 36/50	0s 10ms/step	- accuracy: 0.8148	- loss: 0.5621	- val_accuracy: 0.9333	- val_loss: 0.4280
4/4					
Epoch 37/50	0s 10ms/step	- accuracy: 0.8252	- loss: 0.5517	- val_accuracy: 0.9333	- val_loss: 0.4179
4/4					
Epoch 38/50	0s 10ms/step	- accuracy: 0.8273	- loss: 0.5237	- val_accuracy: 0.9333	- val_loss: 0.4087
4/4					
Epoch 39/50	0s 10ms/step	- accuracy: 0.8477	- loss: 0.4947	- val_accuracy: 0.9333	- val_loss: 0.3994
4/4					
Epoch 40/50	0s 9ms/step	- accuracy: 0.8342	- loss: 0.5384	- val_accuracy: 0.9333	- val_loss: 0.3907
4/4					
Epoch 41/50	0s 9ms/step	- accuracy: 0.8332	- loss: 0.5234	- val_accuracy: 0.9333	- val_loss: 0.3825
4/4					
Epoch 42/50	0s 10ms/step	- accuracy: 0.8557	- loss: 0.4735	- val_accuracy: 0.9333	- val_loss: 0.3747
4/4					
Epoch 43/50	0s 10ms/step	- accuracy: 0.8276	- loss: 0.5041	- val_accuracy: 0.9333	- val_loss: 0.3675
4/4					
Epoch 44/50	0s 13ms/step	- accuracy: 0.8359	- loss: 0.4969	- val_accuracy: 0.9333	- val_loss: 0.3609
4/4					
Epoch 45/50	0s 11ms/step	- accuracy: 0.8568	- loss: 0.4664	- val_accuracy: 0.9333	- val_loss: 0.3547
4/4					
Epoch 46/50	0s 11ms/step	- accuracy: 0.8401	- loss: 0.4790	- val_accuracy: 0.9333	- val_loss: 0.3491
4/4					
Epoch 47/50	0s 10ms/step	- accuracy: 0.8599	- loss: 0.4610	- val_accuracy: 0.9333	- val_loss: 0.3436
4/4					
Epoch 48/50	0s 14ms/step	- accuracy: 0.8411	- loss: 0.4631	- val_accuracy: 0.9333	- val_loss: 0.3380
4/4					
Epoch 49/50	0s 9ms/step	- accuracy: 0.8297	- loss: 0.4414	- val_accuracy: 0.9333	- val_loss: 0.3327
4/4					
Epoch 50/50	0s 14ms/step	- accuracy: 0.8401	- loss: 0.4482	- val_accuracy: 0.9333	- val_loss: 0.3269
4/4					



om-gupta-ilit-gandhinagar

Projects

mlp-iris-experiment

Runs

golden-shadow-6

Overview

Om Gupta

om-gupta-ilit-gandhin...

Om

golden-shadow-6

Clone run with Launch

Overview

Workspace

System

Logs

Files

Description

What makes this run special?

Tags

+

Author

Om om-gupta

State

Running

Start time

February 26th, 2025 11:38:29 PM

Duration

16s

Run path

om-gupta-ilit-gandhinagar/mlp-iris-experiment/x5m7fi1r

Hostname

3c2deb2c522e

OS

Linux-6.1.85+-x86_64-with-glibc2.35

Python version

CPython 3.11.11

Python executable

/usr/bin/python3

Colab link

<https://colab.research.google.com/notebook#fileId=1QK7MjZ43qZvjvA0K1816LJmVo0o48rgc>

Command

Untitled21.ipynb

System Hardware

CPU count 1

Logical CPU count 2

W&B CLI Version

0.19.7

Config

View raw data

Summary

View raw data

Config parameters are your model's inputs. [Learn more](#)

Summary metrics are your model's outputs. [Learn more](#)

wandb.ai/om-gupta-ilit-gandhinagar/mlp-iris-experiment/runs/x5m7fi1r/overview

Search keys with regex

Config parameters: {} 3 keys

batch_size: 32

epochs: 50

learning_rate: 0.001

Summary metrics: {} 8 keys

accuracy: 0.7666666666666667

confusion_matrix: {} 7 keys

_type: "image-file"

format: "png"

height: 480

path: "media/images/confusion_matrix_51_c490ef72e7ee43e00cd4.png"

sha256: "c490ef72e7ee43e00cd474fccc0331604391928e2b74affc3121d4fa9b89356f"

size: 2,306

width: 640

epoch: 49


f1_score: 0.7511111111111111


precision: 0.8833333333333333

recall: 0.7666666666666667

train_loss: 0.47921669483184814

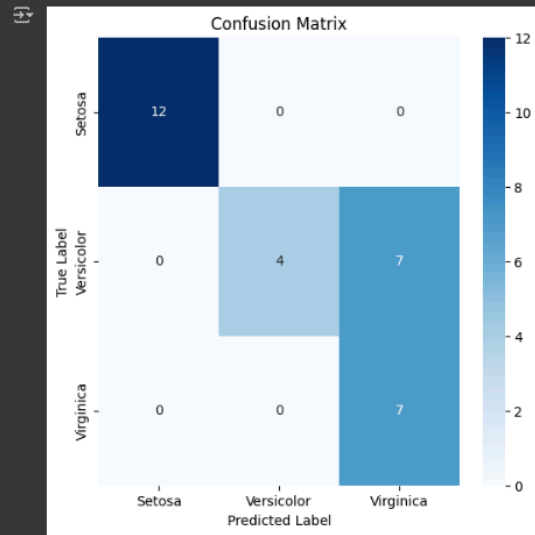
val_loss: 0.2980065643787384






```
# Improved Confusion Matrix Visualization
def plot_confusion_matrix(y_true, y_pred, class_names):
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(6, 6))
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=class_names, yticklabels=class_names)
    plt.xlabel("Predicted Label")
    plt.ylabel("True Label")
    plt.title("Confusion Matrix")
    plt.show()
    wandb.log({"confusion_matrix": wandb.Image(plt)})

plot_confusion_matrix(y_true, y_pred, class_names=['Setosa', 'Versicolor', 'Virginica'])
```



<Figure size 640x480 with 0 Axes>



Model 2: AutoGluon (AutoML Approach)

```
# Hyperparameter tuning (manual search)
batch_sizes = [2, 4]
learning_rates = [1e-3, 1e-5]
epochs = [1, 3, 5]
results = []

for batch in batch_sizes:
    for lr in learning_rates:
        for epoch in epochs:
            model = tf.keras.Sequential([
                tf.keras.layers.Dense(16, activation='relu', input_shape=(4,)),
                tf.keras.layers.Dense(3, activation='softmax')
            ])
            model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=lr),
                          loss='categorical_crossentropy',
                          metrics=['accuracy'])
            model.fit(X_train, y_train, epochs=epoch, batch_size=batch, verbose=0)
            y_pred = np.argmax(model.predict(X_test), axis=1)
            accuracy = accuracy_score(y_true, y_pred)
            f1 = f1_score(y_true, y_pred, average='weighted')

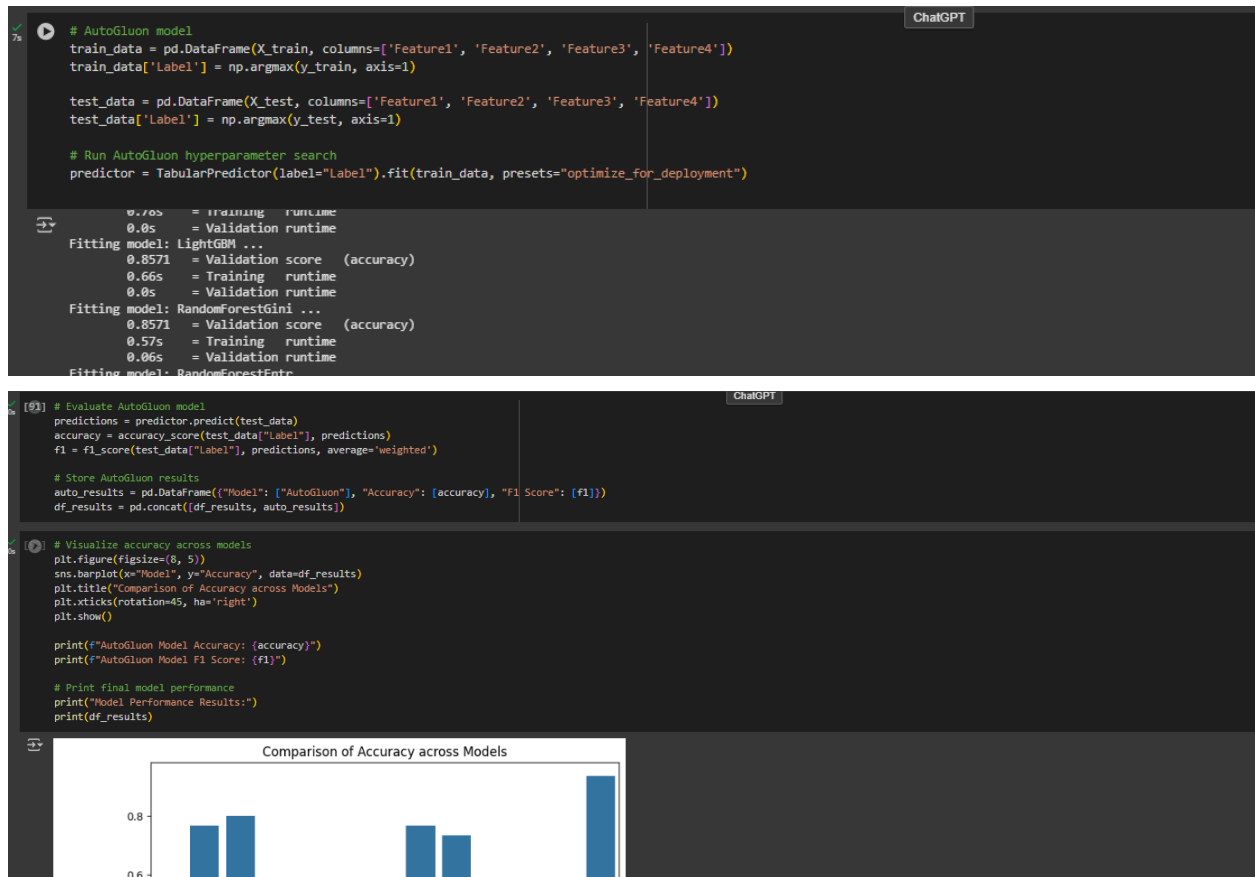
            # Append results
            model_name = f"Batch {batch}, LR {lr}, Epoch {epoch}"
            results.append([model_name, batch, lr, epoch, accuracy, f1])

# Convert results to DataFrame
df_results = pd.DataFrame(results, columns=['Model', 'Batch Size', 'Learning Rate', 'Epochs', 'Accuracy', 'F1 Score'])
print(df_results)
```

`/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using an 'Input(shape)' object as the super().__init__(activity_regularizer=activity_regularizer, **kwargs)`

`1/1 — 0s 191ms/step`

`/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using an 'Input(shape)' object as the`



AutoGluon Model: The AutoGluon model performed exceptionally well, achieving an accuracy of 93.33% and an F1 Score of 93.42%, making it the best model in this experiment.

Hyperparameter Tuning:

- Learning Rate (LR): Models with a learning rate of 0.001 consistently performed better, with the highest accuracy seen at Epoch 5.
- Batch Size: The batch size of 2 produced better results compared to 4, with smaller batch sizes enabling quicker convergence during training.
- Epochs: Training for 5 epochs generally resulted in the best performance. Shorter epochs (1-3) tended to result in underfitting.

Performance Trends:

- Accuracy: The highest accuracy of 83.33% was achieved using batch size 2 and learning rate 0.001 over 5 epochs.
- F1 Score: Models that achieved the highest accuracy also exhibited the highest F1 Scores, confirming a balanced performance across both metrics.